

Introduction to Informa

Assignment Project Exam Help
<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro
Lecture 6: Scoring, Termination and the
Vector Space Model

This lecture; IIR Sections 6.2-6.4.3

- Ranked retrieval
- Scoring documents
- Term frequency
- Collection statistics
- Weighting schemes
- Vector space scoring

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Ranked retrieval

- Thus far, our queries have all been Boolean.
 - Documents either match or don't.
- Good for expert users with precise understanding of their needs a <https://eduassistpro.github.io/>
 - Also good for applications: A can easily consume 1000s of results.
- Not good for the majority of users.
 - Most users incapable of writing Boolean queries (or they are, but they think it's too much work).
 - Most users don't want to wade through 1000s of results.
 - This is particularly true of web search.

Problem with Boolean search: feast or famine

- Boolean queries often result in either too few (=0) or too many (1000s) results.
- Query 1: “standard user dlink 650” → 200,000 hits
- Query 2: “standard user dlink 650 card found”: 0 hits
- It takes a lot of skill to come up with a query that produces a manageable number of hits.
 - AND gives too few; OR gives too many

Ranked retrieval models

- Rather than a set of documents satisfying a query expression, in **ranked retrieval models**, the system returns an ordering over the (top) documents in the collection with
<https://eduassistpro.github.io/>
- **Free text queries**: Rather than a language of operators and expressions, a query is just one or more words in a human language
 Add WeChat edu_assist_pro
- In principle, there are two separate choices here, but in practice, ranked retrieval models have normally been associated with free text queries and vice versa

Feast or famine: not a problem in ranked retrieval

- When a system produces a ranked result set, large result sets are not an issue
 - Indeed, the size of the result set is not an issue
 - We just show it
 - We don't over
- Premise: the ranking algorithm works

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Scoring as the basis of ranked retrieval

- We wish to return in order the documents most likely to be useful to the searcher
- How can we rank-order the documents in the collection with
<https://eduassistpro.github.io/>
- Assign a score to each doc
Add WeChat edu_assist_pro
- This score measures how well the document and query “match”.

Query-document matching scores

- We need a way of assigning a score to a query/document pair
- Let's start with a one-term query
- If the query term is in the document: score should be 0
- The more frequent the query term in the document, the higher the score (should be)
- We will look at a number of alternatives for this.

Boolean Model: Binary term-document incidence matrix

| | Antony and Cleopatra | Julius Caesar | The Tempest | Hamlet | Othello | Macbeth |
|-----------|----------------------|---------------|-------------|--------|---------|---------|
| Antony | 1 | 1 | 0 | 0 | 0 | 1 |
| Brutus | 1 | 1 | 0 | 1 | 0 | 0 |
| Caesar | 1 | | | 1 | 1 | 1 |
| Calpurnia | 0 | | | 0 | 0 | 0 |
| Cleopatra | 1 | 0 | | 0 | 0 | 0 |
| mercy | 1 | 0 | | 1 | 1 | 1 |
| worser | 1 | 0 | 1 | 1 | 1 | 0 |

Each document is represented by a binary vector $\in \{0,1\}^{|V|}$

Take 1: Jaccard coefficient

- Recall from Lecture 3: A commonly used measure of overlap of two sets A and B
- $\text{jaccard}(A, B) = \frac{|A \cap B|}{|A \cup B|}$
- $\text{jaccard}(A, A) = 1$
- $\text{jaccard}(A, B) = 0$ if $A \cap B = \emptyset$
- A and B don't have to be the same size.
- Always assigns a number between 0 and 1.

Jaccard coefficient: Scoring example

- What is the query-document match score that the Jaccard coefficient computes for each of the two documents below?
- Query: *ides of*
- Document 1: *caesar died in*
- Document 2: *the long marc*

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Idea 1: Treat the query as a document

Issues with Jaccard for scoring

- 1 It doesn't consider *term frequency* (how many times a term occurs in a document)
- 2 Rare terms in a collection are more informative than frequent terms. It doesn't consider this information.
<https://eduassistpro.github.io/>
- We need a more sophisticated way of normalizing for length.
Add WeChat edu_assist_pro

Improved Modelling: Term-document count matrices

- Consider the number of occurrences of a term in a document:

- Each document is a count vector in \mathbb{N}^V : a column below

<https://eduassistpro.github.io/>

| | Antony and Cleopatra | Julius Caesar | Hamlet | Othello | Macbeth |
|-----------|----------------------|---------------|--------|---------|---------|
| Antony | 157 | 73 | 0 | 0 | 0 |
| Brutus | 4 | 157 | 0 | 1 | 0 |
| Caesar | 232 | 227 | 0 | 2 | 1 |
| Calpurnia | 0 | 10 | 0 | 0 | 0 |
| Cleopatra | 57 | 0 | 0 | 0 | 0 |
| mercy | 2 | 0 | 3 | 5 | 1 |
| worser | 2 | 0 | 1 | 1 | 0 |

Bag of words model

- Vector representation doesn't consider the ordering of words in a document
- *John is quicker than Mary and Mary is quicker than John* have the <https://eduassistpro.github.io/>
- This is called the bag of words
- In a sense, this is a step back additional index was able to distinguish these two documents.
- We will look at “recovering” positional information later in this course.
- For now: bag of words model

Term frequency tf

- The term frequency $tf_{t,d}$ of term t in document d is defined as the number of times that t occurs in d .
- We want to use tf when computing query-document match scores. <https://eduassistpro.github.io/>
- Raw term frequency is not a good quant:
 - A document with 10 occurrences of a term is more relevant than a document with 1 occurrence of the term.
 - But not 10 times more relevant.
- Relevance does not increase proportionally with term frequency.

Idea 2: Normalized tf is an important scoring factor

Log-frequency weighting

- The log frequency weight of term t in d is

$$w_{t,d} = \begin{cases} 1 + \log_{10} \text{tf}_{t,d}, & \text{if } \text{tf}_{t,d} > 0 \\ 0, & \text{otherwise} \end{cases}$$

- $0 \rightarrow 0, 1 \rightarrow 1, \dots, 4 \rightarrow 4$, etc.
- Score for a document d over terms t in both q and d :
- $$\text{score} = \sum_{t \in q \cap d} (1 + \log \text{tf}_{t,d})$$
- The score is 0 if none of the query terms is present in the document.

Document frequency

- Rare terms are more informative than frequent terms
 - Recall stop words
- Consider a term *arachn* are in the collection (e.g. <https://eduassistpro.github.io/>)
- A document containing this term is likely to be relevant to the query *arachn*
- → We want a high weight for rare terms like *arachnocentric*.

Document frequency, continued

- Frequent terms are less informative than rare terms
- Consider a query term that is frequent in the collection (e.g., *high*, *increase*, *line*)
- A document containing this term is more likely to be relevant than a document that doesn't
- But it's not a sure indicator
- → For frequent terms, we want high positive weights for words like *high*, *increase*, and *line*
- But lower weights than for rare terms.
- We will use document frequency (df) to capture this.

idf weight

- df_t is the document frequency of t : the number of documents that contain t
 - df_t is an inverse measure of the informativeness of t
 - $df_t \leq N$
- We define the t frequency) of t by

$$idf_t = \log_{10} (N/df_t)$$

- We use $\log (N/df_t)$ instead of N/df_t to “dampen” the effect of idf.

Idea 3: Normalized idf is an important scoring factor

Will turn out the base of the log is immaterial.

idf example, suppose $N = 1$ million

| term | df_t | idf_t |
|-----------|-----------|---------|
| calpurnia | 1 | |
| animal | 100 | |
| sunday | | |
| fly | | |
| under | | |
| the | 1,000,000 | |

$$idf_t = \log_{10} (N/df_t)$$

There is one idf value for each term t in a collection.

Effect of idf on ranking

- Does idf have an effect on ranking for one-term queries, like
 - iPhone
- idf has no effect on ranking for one-term queries
 - idf affects the ranking of doc queries with at least two terms
 - For the query *capricious person*, idf weighting makes occurrences of *capricious* count for much more in the final document ranking than occurrences of *person*.

Collection vs. Document frequency

- The collection frequency of t is the number of occurrences of t in the collection, counting multiple occurrences.

- Example: <https://eduassistpro.github.io/>

| Word | Collection frequency | Document frequency |
|------------------|----------------------|--------------------|
| <i>insurance</i> | 10440 | 3997 |
| <i>try</i> | 10422 | 8760 |

- Which word is a better search term (and should get a higher weight)?

tf-idf weighting

- The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$w_{t,d} = \left(1 + \frac{N}{df_t}\right)$$

Assignment Project Exam Help
<https://eduassistpro.github.io/>

- Best known weighting scheme for information retrieval
 - Note: the “-” in tf-idf is a hyphen, not a minus sign!
 - Alternative names: tf.idf, tf x idf
- Increases with the number of occurrences within a document
- Increases with the rarity of the term in the collection

Binary \rightarrow count \rightarrow weight matrix

| | Antony and Cleopatra | Julius Caesar | The Tempest | Hamlet | Othello | Macbeth |
|-----------|----------------------|---------------|-------------|--------|---------|---------|
| Antony | 5.25 | 3.18 | 0 | 0 | 0 | 0.35 |
| Brutus | 1.21 | 6.1 | 0 | 1 | 0 | 0 |
| Caesar | 8.59 | | | 1.51 | 0.25 | 0 |
| Calpurnia | 0 | | | 0 | 0 | 0 |
| Cleopatra | 2.85 | 0 | | 0 | 0 | 0 |
| mercy | 1.51 | 0 | | 0.12 | 5.25 | 0.88 |
| worser | 1.37 | 0 | 0.11 | 4.15 | 0.25 | 1.95 |

Each document is now represented by a real-valued vector of tf-idf weights $\in \mathbb{R}^{|V|}$

Documents as vectors

- So we have a $|V|$ -dimensional vector space
- Terms are axes of the space
- Documents are points in this space
- Very high-dimensional vectors
dimensions when you apply web search engine
- These are very sparse vectors - most entries are zero.

Queries as vectors

- Key idea 1: Do the same for queries: represent them as vectors in the space
- Key idea 2: Rank documents according to their proximity to t
- proximity = similarity of ve
- proximity \approx inverse of dista
- Recall: We do this because we want to get away from the you're-either-in-or-out Boolean model.
- Instead: rank more relevant documents higher than less relevant documents

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Formalizing vector space proximity

- First cut: distance between two points
 - (= distance between the end points of the two vectors)
- Euclidean distance?
- Euclidean distance <https://eduassistpro.github.io/>
- ... because Euclidean distance for vectors of different lengths.

Why distance is a bad idea

The Euclidean
distance between q

and d_2 is large even

though the

distribution of ter

in the query q and the
distribution of

terms in the

document d_2 are

very similar.

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Use angle instead of distance

- Thought experiment: take a document d and append it to itself. Call this document d' .
- “Semantically” d and d' have the same content
- The Euclidean distance between two documents can be quite large
- The angle between the two documents is 0, corresponding to maximal similarity.
- Key idea: Rank documents according to angle with query.

From angles to cosines

- The following two notions are equivalent.
 - Rank documents in decreasing order of the angle between query and document
 - Rank documents in decreasing order of $\cos(\text{angle}(\text{query}, \text{document}))$
- Cosine is a monotonically decreasing function for the interval $[0^\circ, 180^\circ]$

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

From angles to cosines

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

- But how – *and why* – should we be computing cosines?

Length normalization

- A vector can be (length-) normalized by dividing each of its components by its length – for this we use the L_2 norm: **Assignment Project Exam Help**
 $\frac{\vec{v}}{\|\vec{v}\|_2}$
<https://eduassistpro.github.io/>
- Dividing a vector by its L_2 norm (length) vector (on surface of unit sphere)
 Add WeChat edu_assist_pro
- Effect on the two documents d and d' (d appended to itself) from earlier slide: they have identical vectors after length-normalization.
 - Long and short documents now have comparable weights

cosine(query,document)

Dot product Unit vectors

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \bullet \vec{d}}{\|\vec{q}\| \|\vec{d}\|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

Assignment Project Exam Help
<https://eduassistpro.github.io/>
 Add WeChat edu_assist_pro

q_i is the tf-idf weight of term i in the query

d_i is the tf-idf weight of term i in the document

$\cos(\vec{q}, \vec{d})$ is the cosine similarity of \vec{q} and \vec{d} ... or,
 equivalently, the cosine of the angle between \vec{q} and \vec{d} .

Cosine for length-normalized vectors

- For length-normalized vectors, cosine similarity is simply the dot product (or scalar product):

Assignment Project Exam Help

$$\cos(\vec{q}, \vec{d}_i)$$

<https://eduassistpro.github.io/>
Add WeChat edu_assist_pro

for q, d length-normalized.

Cosine similarity illustrated

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Cosine similarity amongst 3 documents

How similar are
the novels

SaS: *Sense and*

Sensibility

PaP: *Pride and*

Prejudice, and

WH: *Wuthering*

Heights?

| term | SaS | PaP | WH |
|-----------|-----|-----|----|
| affection | 115 | 58 | 20 |
| gossip | 7 | 11 | 6 |
| wuthering | 0 | 0 | 38 |

Term frequencies (counts)

Note: To simplify this example, we don't do idf weighting in this example.

3 documents example contd.

Log frequency weighting

After length normalization

| term | SaS | PaP | WH | term | SaS | PaP | WH |
|-----------|------|-----|------|------|-------|-------|-------|
| affection | 3.06 | | | | 0.789 | 0.832 | 0.524 |
| jealous | 2.00 | | | | 0.515 | 0.555 | 0.465 |
| gossip | 1.30 | 0 | 1.78 | g | .335 | 0 | 0.405 |
| wuthering | 0 | 0 | 2.58 | | 0 | 0 | 0.588 |

$\cos(\text{SaS}, \text{PaP}) \approx$

$$0.789 \times 0.832 + 0.515 \times 0.555 + 0.335 \times 0.0 + 0.0 \times 0.0 \\ \approx 0.94$$

$\cos(\text{SaS}, \text{WH}) \approx 0.79$

$\cos(\text{PaP}, \text{WH}) \approx 0.69$

Why do we have $\cos(\text{SaS}, \text{PaP}) > \cos(\text{SaS}, \text{WH})$?

Computing cosine scores

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

tf-idf weighting has many variants

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Columns headed 'n' are acronyms for weight schemes.

Why is the base of the log in idf immaterial?

Weighting may differ in queries vs documents

- Many search engines allow for different weightings for queries vs. documents
- SMART Notation: denotes the combination in use in an engine, with using the acronyms from <https://eduassistpro.github.io/>
- A very standard weighting
 - Document: logarithmic tf (l as first character), no idf and cosine normalization
 - Query: logarithmic tf (l in leftmost column), idf (t in second column), cosine normalization ...

A bad idea?

tf-idf example: Inc.Itc

Document: *car insurance auto insurance*

Query: *best car insurance*

| Term | Query | | | | | | Document | | | | Pro d |
|-----------|--------|-------|-------|--------|--------|-------|----------|--------|--------|------|----------|
| | tf-raw | tf-wt | idf | tf-idf | tf-raw | tf-wt | idf | tf-idf | n'lize | | |
| auto | 0 | 0 | 50000 | 0 | 1 | 1 | 1 | 0.52 | 0 | 0 | |
| best | 1 | 1 | 50000 | 1.3 | 0 | 0 | 0 | 0 | 0 | 0 | |
| car | 1 | 1 | 10000 | 2.0 | 1 | 1 | 1 | 0.52 | 0.52 | 0.27 | |
| insurance | 1 | 1 | 1000 | 3.0 | 2 | 1.3 | 1.3 | 0.68 | 0.53 | | |

Exercise: what is N , the number of docs?

$$\text{Doc length} = \sqrt{1^2 + 0^2 + 1^2 + 1.3^2} \approx 1.92$$

$$\text{Score} = 0 + 0 + 0.27 + 0.53 = 0.8$$

Representation/feature perspective

- Inc.ltc
 - doc vector:
 - tf-vector
 - normalized
 - query vector
 - tf-idf-vector
 - normalized to unit length
 - score = similarity = inner product between the two vectors

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

Summary – vector space ranking

- Represent the query as a weighted tf-idf vector
- Represent each document as a weighted tf-idf vector
- Compute the cosine similarity for the query vector and each document vector
- Rank documents with respect to the query by score
- Return the top K (e.g., $K = 10$) to the user

Resources for today's lecture

- IIR 6.2 – 6.4.3
- <http://www.assignment-project-exam-help.com/n-retrieval-tutorial/cosine-similarity-tutorial/>
<https://github.com/eduassistpro>
 - Term weighting and cosine similarity tutorial for SEO folk!

Add WeChat edu_assist_pro